



UNIVERSITÀ
di **VERONA**

Dipartimento
di **INFORMATICA**



Part I

Laurea magistrale in Ingegneria e Scienze Informatiche
Laurea Magistrale in Medical Bioinformatics

Nicola Bombieri – Federico Busato

Agenda

- What is CUDA?
- Calling a Device Function
- CUDA Concepts
- CUDA Keywords:
 - Built-I
- Memory Management
- Error handling
- CUDA Compiler Driver
- Timing the code
- Examples
- Exercises
- References
- Books

What is CUDA?

- **CUDA** (**C**ompute **U**nified **D**evice **A**rchitecture) is a parallel computing platform and programming model that makes using a GPU for general purpose computing
- CUDA is small extension of C/C++ language
- With CUDA you can accelerate your C/C++ code by moving the computationally intensive portions of your code to an NVIDIA GPU

Calling a Kernel Function

- A kernel is a function callable from the host and executed on the CUDA device
- The kernel is the heart of your CUDA code

```
__global__ void MyKernel (int* Array, int size, ...) {  
    ...  
}  
  
int main() {  
    ...  
    MyKernel<<<...>>>(Array, size, ...);  
    ...  
}
```

- A kernel is defined using the __global__ declaration specifier
- The number of CUDA threads that execute a kernel is specified using a new <<<...>>> execution configuration syntax
- The kernel must have return type `void`

CUDA Architecture Overview



Built-in Variables

- **threadIdx.x**

Contains the thread index within the block (dimension X)

- **blockIdx.x**

Contains the thread index within the grid (dimension X)

- **blockDim.x**

Contains the dimension of the block (number of threads) (dimension X)

- **gridDim.x**

Contains the dimension of the grid (number of blocks) (dimension X)

- ✓ Common case of threads indexing within the Kernel:

GlobalThreadIndex = **blockIdx.x** * **blockDim.x** + **threadIdx.x**

- ✓ Useful Kernel Function:

printf(...) // need synchronization (implicit/explicit) after kernel

Common CUDA Code Steps

- 1) Initialize host data structures
- 2) Allocate device data structures `cudaMalloc`
- 3) Copy host data structures to device `cudaMemcpy`
- 4) Invoke kernel function `MyKernel<<<>>>`
- 5) Copy device result to host `cudaMemcpy`
- 6) Free host and device data structures `cudaFree`

When the process terminates all allocated memory is released

- 7) Reset the device
`cudaDeviceReset`

CUDA memory management (I)

Device Variables (symbols):

```
cudaMemcpyToSymbol(symbol, void* src, size_t size, offset=0)
```

- Copies SIZE bytes from the memory area pointed by SRC to the memory area pointed to by OFFSET bytes from the start of SYMBOL
- Pointer arithmetic is not valid on SYMBOL

```
cudaMemcpyFromSymbol(void* dest, symbol, size_t size, offset=0)
```

- Copies SIZE bytes from the memory area OFFSET bytes from the start of SYMBOL to the memory area pointed to by DEST
- Pointer arithmetic is not valid on SYMBOL

CUDA memory management

Example (I)

```
__global__ void MyKernel (int* Array, int size, ...) { ... }

int main() {
    ...
    int* devArray;
    int size = 120;
    cudaMalloc(&devArray, size * sizeof(int));
    cudaMemcpy(devArray, hostArrayA, size * sizeof(int), cudaMemcpyHostToDevice);
    //valid: cudaMemcpy(devArray + 4, hostArrayA, ...);
    MyKernel<<<...>>>(devArray, size, ...);
    cudaMemcpy(hostArrayB, devArray, size * sizeof(int), cudaMemcpyDeviceToHost);
    cudaFree(devArray);
    cudaDeviceReset();
}
```

CUDA memory management

Example (II)

```
__device__ sArray[128];

__global__ void MyKernel (...) {
    ...
    sArray[4] = 14;
}

int main() {
    ...
    cudaMemcpyToSymbol(sArray, hostArrayA, 4 * sizeof(int), 4);
    MyKernel<<<...>>>(...);
    cudaMemcpyFromSymbol(hostArrayB, sArray, 4 * sizeof(int), 4);
}
```

CUDA Error Handling (I)


- Every CUDA call (except kernel launches) return an error code of type `cudaError_t`
 - No error = "`cudaSuccess`"
 - otherwise an error code

```
cudaError_t err = cudaMalloc( &fooPtr, -1 );  
if ( cudaSuccess != err )  
    printf("Error! : %d\n", err);
```

- CUDA kernel invocations do not return any value. Error from a CUDA kernel call can be checked after its execution by calling `cudaGetLastError()`

```
fooKernel<<< x, y >>>(); // Kernel call  
cudaDeviceSynchronize();  
cudaError_t err = cudaGetLastError();  
. . .
```

Important!!: Need to
Synchronize Host and
Device



- A human-readable description of the error can be obtained from
`char *cudaGetErrorString(cudaError_t code);`

CUDA Error Handling (II)

Elegant solutions:

```
#define cudaError( msg )    getLastCudaError(msg, __FILE__, __LINE__)
#define CudaSafeCall( err ) cudaSafeCall(err, __FILE__, __LINE__ )

inline void getLastCudaError(const char *errorMessage,
                             const char *file, const int line) {
    cudaDeviceSynchronize();
    cudaError_t err = cudaGetLastError();
    if (cudaSuccess != err) {
        std::cerr << file << " (" << line << ") : getLastCudaError()
        CUDA error : " << errorMessage << " : (" << (int) err << " ) "
        << cudaGetErrorString(err) << std::endl << std::endl;

        std::exit(EXIT_FAILURE);
    }
}

inline void cudaSafeCall(cudaError err, const char *file,
                          const int line) {
    __getLastCudaError("", file, line);
}
```

Alternative timing of your Code (I)

Events:

```
cudaEvent_t startTimeCuda, stopTimeCuda;  
cudaEventCreate(&startTimeCuda);  
cudaEventCreate(&stopTimeCuda);  
  
cudaEventRecord(startTimeCuda, 0); //0 is the default stream  
  
...code...           //also host code and cuda functions  
Kernel<<<>>>(); //one or more  
...code...  
  
cudaEventRecord(stopTimeCuda, 0);  
cudaEventSynchronize(stopTimeCuda);  
float msTime;  
cudaEventElapsedTime(&msTime, startTimeCuda, stopTimeCuda);
```

- Compute the elapsed time in milliseconds between two recorded events

Alternative timing of your Code (II)

Host Timer

```
// host timer init
// host timer start

...code...
Kernel<<<>>>();
...code...

cudaDeviceSynchronize();
// host timer stop
```

- Explicit synchronization barrier cudaDeviceSynchronize() is required to block CPU execution until all previously issued commands on the device have completed
- Without this barrier, this code would measure the kernel launch time and not the kernel execution time

IMPORTANT: the execution is non-deterministic -> Sometimes it is better to repeat many times the computation and then compute the average time to get a better approximation of the execution time.

CUDA Compiler Driver (I)

- Any source file containing CUDA language extensions must be compiled with NVCC
- NVCC is a compiler driver
 - It automatically invokes all the necessary tools and compilers like `cudacc`, `g++`, ...
- Any executable with CUDA code requires the CUDA runtime library (`cudart`)
- This library must be in the standard library path or the environment variable `LD_LIBRARY_PATH` must contain the path to this library
 - Common case:

```
export LD_LIBRARY_PATH=/usr/local/cuda/<lib/lib64>:$LD_LIBRARY_PATH
export PATH=/usr/local/cuda/bin:$PATH
```

- To compile a CUDA program:

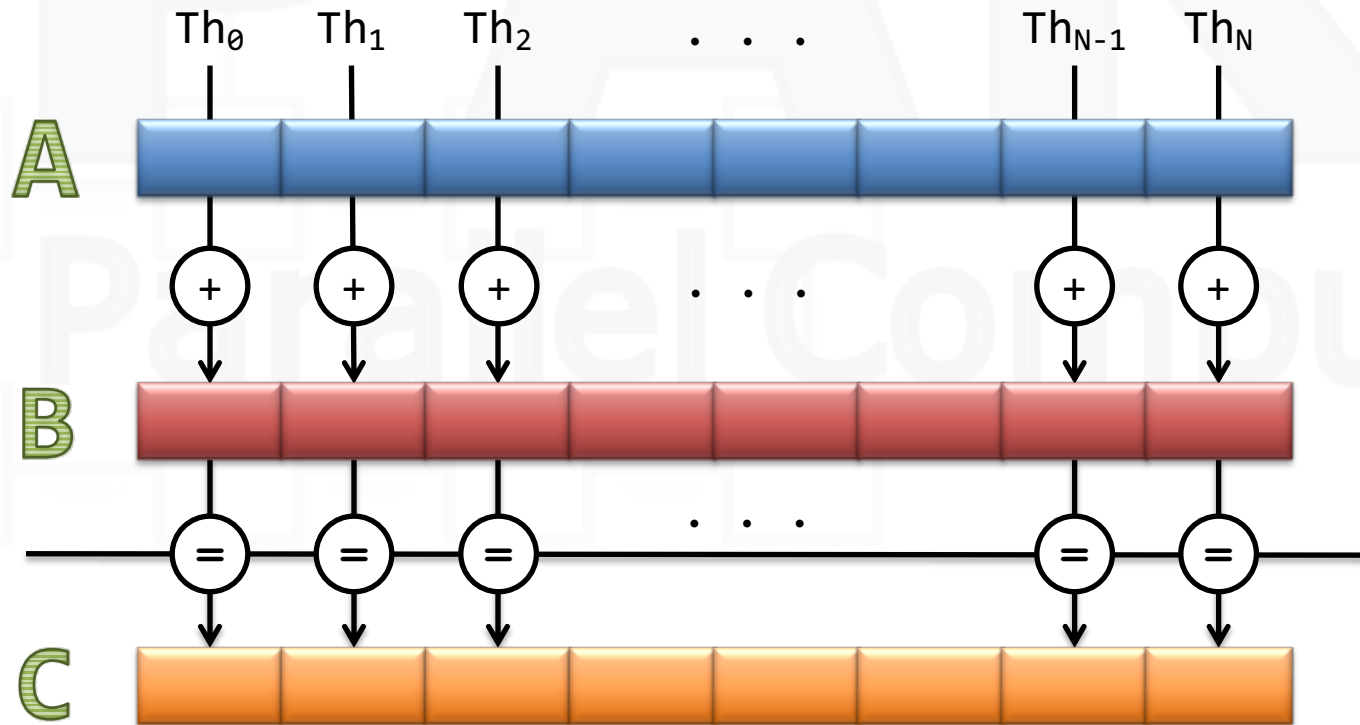
```
nvcc source.cu -o out.x
```

Some parameters:

```
-std=c++11    -arch=sm_62
```


Examples

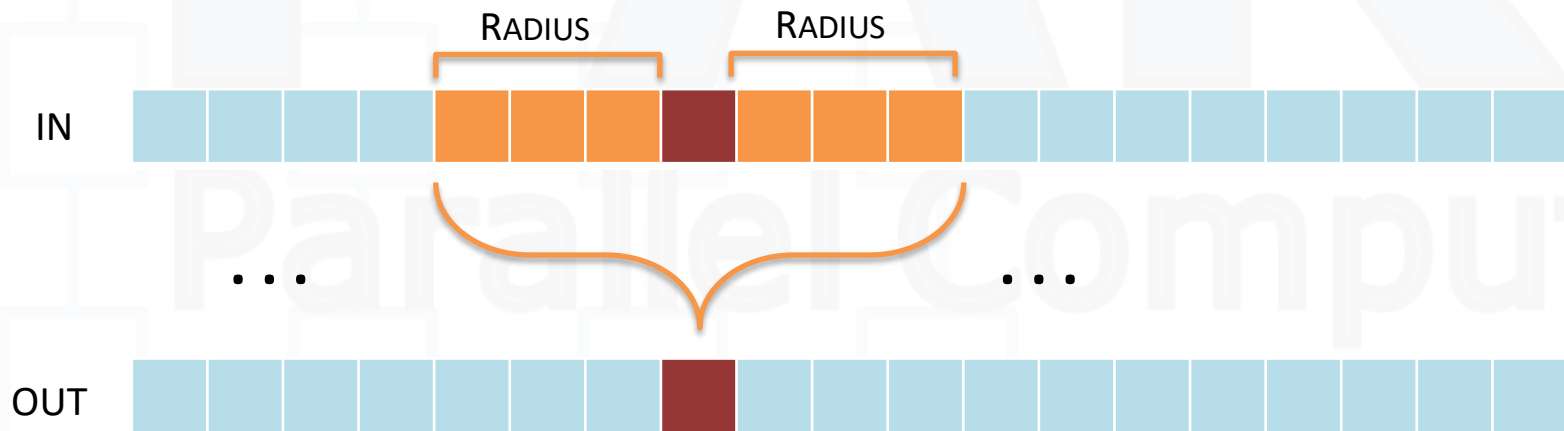
- Hello World
- Vector addition:



Exercise

1D STENCIL (I)

- Consider applying a 1D stencil to a 1D array of elements
 - Each output element is the sum of input elements within a **radius**
 - If radius is 3, then each output element is the sum of 7 input elements




- Fundamental to many algorithms Standard discretization methods, interpolation, convolution, filtering

Exercise

1D STENCIL (II)

- Each thread processes one output element
- With radius 3, each input element is read seven times
 - Input elements are read several times

THREAD COMPUTATION



NOTE: handle left and right bounds



Exercise

MATRIX MULTIPLICATION

Simplest Version:

- The matrix must be linearized:

$$A[i][j] = A[i * \text{rows} + j]$$

- Set up 2D Grid and 2D Blocks
- Each thread takes one row and one column
- Block size divide N, for each dimensions
- Requires N x N threads. e.g. N = 100 -> 10'000 Threads

What is the speedup against the sequential version for large values?

0	1	2
3	4	5
6	7	8

A



0	1	2
3	4	5
6	7	8

B




0	1	2
3	4	5
6	7	8

C

Exercises

MATRIX TRANSPOSE



Simplest Version:

- Each thread takes one cell
- The kernel produce on output the transpose matrix
- Requires $N \times N$ threads. e.g. $N = 100 \rightarrow 10'000$ Threads

References

- CUDA C Programming Guide

<http://docs.nvidia.com/cuda/cuda-c-programming-guide/>

- CUDA Runtime API

<http://docs.nvidia.com/cuda/cuda-runtime-api/>







http://developer.download.nvidia.com/compute/cuda/4_1/rel/toolkit/docs/online/modules.html

(Memory Management)

- CUDA Math API

<http://docs.nvidia.com/cuda/cuda-math-api/>

CUDA Programming Books



CUDA Advanced Books

